

Joseph Le

Palo Alto, CA – josephle@stanford.edu – github.com/JosephLeKH – josephle.dev

EDUCATION

Stanford University

Palo Alto, CA

Bachelor of Science in Computer Science (Systems and AI Track), GPA: 3.97/4.0

June 2027

- Relevant Coursework: Data Structures & Algorithms, Operating Systems, Computer Systems, Artificial Intelligence, Machine Learning, Databases, Natural Language Processing, Linear Algebra
- Leadership & Activities: VP, CS + Social Good; Teaching Assistant, CS 106B; Stanford AI Club

WORK EXPERIENCE

Member of Technical Staff

Palo Alto, CA

Vimes

Feb 2026 - Present

- Built a repeatable data-ingestion pipeline (AWS ECS Fargate, RDS) unifying records from multiple independent agency systems into one searchable store, normalizing heterogeneous schemas into a common document model.
- Re-architected the platform's authorization layer into a cache-backed model with atomic writes and automated migration tooling, hardening access control across every department without disrupting existing users.
- Engineered an AI-powered form-completion backend using context-aware LLM prompting, cutting token usage while improving suggestion accuracy and keeping sensitive data scoped server-side.

Software Engineering Intern

Cupertino, CA

Apple

Jun 2026 - Sep 2026

- Building a Swift iOS application for Apple's internal engineering Expo, implementing interactive venue mapping, personalized calendar management, real-time push notifications, and AI-powered session recommendations for a large-scale internal audience.

Software Engineer

Palo Alto, CA

Stanford Carta Lab

Dec 2025 – June 2026

- Built core planning features for Stanford's four-year course planner, implementing interactive scheduling workflows, prerequisite and unit constraint validation, and robust client-server state synchronization across complex multi-quarter plans.
- Identified and resolved 20+ frontend state and edge-case bugs, improving UI consistency and reducing student-reported planning errors in a high-traffic production environment.

Machine Learning Research Intern

Palo Alto, CA

Neurotrack

Oct 2025 – Jun 2026

- Developed and evaluated machine learning models to detect and predict Mild Cognitive Impairment (MCI) and dementia from large-scale cognitive assessment datasets.
- Designed scalable data preprocessing, feature engineering, and validation pipelines to support experimentation and clinical research.

Founding Engineer

Palo Alto, CA

Bases Insight

Aug 2025 - Mar 2026

- Architected and built a production financial analysis platform powering automated valuation workflows (DCF, comps, three-statement modeling) with an asynchronous FastAPI backend, React frontend, and PostgreSQL database.
- Designed distributed async processing pipelines for document extraction and validation, improving model generation throughput and system scalability.
- Implemented CI/CD pipelines with GitHub Actions and automated testing, reducing friction and increasing release reliability.

Data Science Intern

Austin, TX

Bonterra

Jun 2025 - Aug 2025

- Developed predictive models using scikit-learn, XGBoost, and CausalML to estimate treatment effects across 2.7M+ records.
- Designed and optimized large ETL pipelines with SQL, Pandas, and NumPy for feature engineering, model training, and evaluation.

TECHNICAL PROJECTS

SemEval-2025: Multilingual Emotion Prediction — aclanthology.org/2025.semeval-1.244/

- Fine-tuned multilingual BERT and XLM-RoBERTa models for emotion intensity regression across 11 languages, implementing custom training pipelines, cross-lingual evaluation, and metric optimization under benchmark constraints.

Wildfire Insurance Analysis — Python, Pandas, SQL, Regression

- Built end-to-end data pipelines and regression models to analyze California wildfire severity (2016–2023), applying time-series analysis and statistical modeling to quantify insurance non-renewal trends and FAIR Plan dependence.

TECHNICAL SKILLS

Languages: Python, Java, C++, C, Swift, JavaScript, TypeScript, SQL, Go, Bash

Frameworks & APIs: React, React Native, Node.js, Express, Spring, FastAPI, Flask, REST APIs, GraphQL

Databases & Storage: PostgreSQL, MySQL, MongoDB, Redis, Elasticsearch, Database Design, Query Optimization

Systems & Cloud: AWS (EC2, S3, Lambda, IAM), Docker, Kubernetes, Microservices, Distributed Systems, Scalability, Load Balancing, Caching, CI/CD, GitHub Actions, Unix/Linux, Networking, Containerization

Machine Learning: PyTorch, TensorFlow, scikit-learn, XGBoost, HuggingFace Transformers, Pandas, NumPy, NLP, LLM Integration